*Note: replace "**$USER**" with your username

Setup Steps:

- Create new env "climategpt3" in */nobackup*
  - `conda create --prefix=/nobackup/users/$USER/.conda/envs/climategpt3`
- Create jupyter kernel:
  - `python -m ipykernel install --user --name climategpt3`
- Add Conda Channel:
  - `conda config --prepend channels \ https://public.dhe.ibm.com/ibmdl/export/pub/software/server/ibm-ai/conda/`
  - `conda config --prepend channels \ https://opence.mit.edu`

Getting pytorch:

- `module load anaconda3`
- `srun --gres=gpu:1 -n 4 --time 1:00:00 --pty /bin/bash`
- `module load anaconda3`
- `module load cuda/11.4`
- `source activate /nobackup/users/$USER/.conda/envs/climategpt3`
- `conda install pytorch=1.12.1=cuda11.4_py39_1`
  - This automatically installs compatible cudnn & cudatoolkit versions
  - *note: ensure there is enough space in the home directory – for some reason, installing pytorch used up 9GB in /home in addition to 35GB in /nobackup
    - du -sh /nobackup/users/**$USER**
    - du -sh /home/**$USER**

Getting transformers:

- `module load anaconda3`
- `source activate /nobackup/users/$USER/.conda/envs/climategpt3`
- `conda install transformers=4.32.1`

Afterward on Jupyter notebook:

- Activate climategpt3 kernel

Troubleshooting:

- For disk space-related errors (especially if the home directory is almost full or for "disk-quota exceeded" error), try:
  - `pip cache purge`
  - `conda clean --all`
- `RuntimeError: Failed to import transformers.models.llama.tokenization_llama_fast because of the following error (look up to see its traceback): tokenizers>=0.13.3 is required for a normal functioning of this module, but found tokenizers==0.11.4.`
  - `conda install tokenizers==0.13.3`